

# Network IPC: Sockets

---

Advanced Programming in the UNIX Programming Environment

Chun-Ying Huang <chuang@cs.nctu.edu.tw>

# Introduction

---

The classical IPC's allow processes running on the same computer to communicate with one another

To allow processes running on different computers with one another:  
Use Network IPC

Network IPC can be used for both inter-machine communication and intra-machine communication

In this Chapter, we focus majorly on TCP/IP sockets

# Socket Descriptors

---

A socket is an abstraction of a communication endpoint

Socket descriptors are implemented as file descriptors in the UNIX System

Many of the functions that deal with file descriptors, such as *read* and *write*, will work with a socket descriptor

# Socket Descriptors (Cont'd)

---

## Creating a socket

- `int socket(int domain, int type, int protocol);`
- Returns: file (socket) descriptor if success, or -1 on error

domain: `AF_INET`, `AF_INET6`, `AF_UNIX`, `AF_UNSPEC`

type: `SOCK_DGRAM`, `SOCK_RAW`, `SOCK_STREAM`

protocol: 0 (default), `IPPROTO_TCP`, `IPPROTO_UDP`, ...

Note: the above constants (defines) may locate in different header files

- Most of domain/type resides in `sys/socket.h`
- `IPPROTO_*` resides in `netinet/in.h`

# Socket Descriptors and File I/O Functions

Function	Behavior with socket
<code>close</code>	deallocates the socket
<code>dup, dup2</code>	duplicates the file descriptor as normal
<code>fchdir</code>	fails with <code>errno</code> set to <code>ENOTDIR</code>
<code>fchmod</code>	unspecified
<code>fchown</code>	implementation defined
<code>fcntl</code>	some commands supported, including <code>F_DUPFD</code> , <code>F_GETFD</code> , <code>F_GETFL</code> , <code>F_GETOWN</code> , <code>F_SETFD</code> , <code>F_SETFL</code> , and <code>F_SETOWN</code>
<code>fdatasync, fsync</code>	implementation defined
<code>fstat</code>	some <code>stat</code> structure members supported, but how left up to the implementation
<code>ftruncate</code>	unspecified
<code>ioctl</code>	some commands work, depending on underlying device driver
<code>lseek</code>	implementation defined (usually fails with <code>errno</code> set to <code>ESPIPE</code> )
<code>read</code>	equivalent to <code>recv</code> without any flags
<code>write</code>	equivalent to <code>send</code> without any flags

# Release a Socket Descriptor

---

Communication on a socket is bidirectional

We can disable I/O on a socket with the shutdown function

## Synopsis

- `int shutdown(int sockfd, int how);`
- Returns: zero if success, or -1 on error

how: SHUT\_RD, SHUT\_WR, and SHUT\_RDWR

## Why do we need shutdown?

- shutdown closes the socket descriptor immediately (independent of the number of references to the descriptor)
- shutdown is able to half-close a socket descriptor

# Addressing

---

How to identify a socket?

We need addressing schemes

- AF\_UNIX
  - a pathname
- AF\_INET + SOCK\_STREAM + IP\_PROTO\_TCP
  - IPv4 address and TCP port number
- AF\_INET + SOCKET\_DGRAM + IP\_PROTO\_UDP
  - IPv4 address and UDP port number
- AF\_INET6 + SOCK\_STREAM + IP\_PROTO\_TCP
  - IPv6 address and TCP port number
- AF\_INET6 + SOCKET\_DGRAM + IP\_PROTO\_UDP
  - IPv6 address and UDP port number

# Address and Byte Ordering



---

When communicating with processes running on the same computer, we generally don't have to worry about byte ordering

It is a problem when addresses are represented in numeric forms

The byte order is a characteristic of the processor architecture

- It only affects larger data types, e.g., an integer





# Address and Byte Ordering (Cont'd)

---

The byte ordering becomes visible to applications when they exchange formatted data

The TCP/IP protocol suite uses big-endian byte order

- The network byte order

Four common functions are provided to convert byte orders

- `uint32_t htonl(uint32_t hostint32);`
- `uint16_t htons(uint16_t hostint16);`
- `uint32_t ntohl(uint32_t netint32);`
- `uint16_t ntohs(uint16_t netint16);`

# Address Formats: Generic

---

The generic sockaddr structure on Linux

```
struct sockaddr {  
    sa_family_t sa_family;    /* address family */  
    char sa_data[14];        /* variable-length address */  
};
```

However, it may be a little bit different on some other systems

```
struct sockaddr {  
    unsigned char sa_len      /* total length */  
    sa_family_t sa_family;    /* address family */  
    char sa_data[14];        /* variable-length address */  
};
```

# Address Formats: IPv4 on Linux

---

```
typedef uint16_t in_port_t;
typedef uint32_t in_addr_t;

struct in_addr {
    in_addr_t s_addr;           /* IPv4 address */
};

struct sockaddr_in {
    sa_family_t sin_family;     /* address family */
    in_port_t sin_port;         /* port number */
    struct in_addr sin_addr;     /* IPv4 address */
    unsigned char sin_zero[8];
};
```

# Address Formats: IPv6 on Linux

---

```
typedef uint16_t in_port_t;

struct in6_addr {
    union {                                /* IPv4 address */
        uint8_t  __u6_addr8 [16];
        uint16_t __u6_addr16[8];
        uint32_t __u6_addr32[4];
    } __in6_u;
};

struct sockaddr_in6 {
    sa_family_t sin_family;    /* address family */
    in_port_t sin6_port;      /* port number */
    uint32_t sin6_flowinfo;   /* IPv6 flow info */
    struct in6_addr sin6_addr; /* IPv6 address */
    uint32_t sin6_scope_id;    /* IPv6 scope id */
};
```

# Conversion of Address

---

To print an address in a readable format, we have to convert a numeric address to a text

- `const char *inet_ntop(int af, const void *src, char *dst, socklen_t cnt);`
- Returns: pointer to address string on success, or NULL on error

To convert a text formatted address to a numeric address

- `int inet_pton(int af, const char *src, void *dst);`
- Returns: 1 on success, 0 if the format is invalid, or -1 on error and `errno=EAFNOSUPPORT`

# Example: Address Conversion

---

See `netipc/addr.c`

```
if(inet_pton(AF_INET, argv[1], &addr4) == 1) {  
    printf("IPv4: 0x%08x\n", htonl(addr4.s_addr));  
}  
  
if(inet_pton(AF_INET6, argv[1], &addr6) == 1) {  
    printf("IPv6: 0x%08x%08x%08x%08x\n",  
        htonl(addr6.__in6_u.__u6_addr32[0]),  
        htonl(addr6.__in6_u.__u6_addr32[1]),  
        htonl(addr6.__in6_u.__u6_addr32[2]),  
        htonl(addr6.__in6_u.__u6_addr32[3]));  
}
```

# Address Lookup – Known Hosts

---

How do we get all known hosts in the system?

Check the /etc/hosts file

In a UNIX program, we can get all known hosts by the function:

- `struct hostent *gethostent(void);`
- Returns: valid pointer if success, or NULL on error
- NOTE: `gethostent()` is not thread-safe

```
struct hostent {  
    char  *h_name;           /* official name of host */  
    char **h_aliases;        /* alias list */  
    int    h_addrtype;       /* host address type */  
    int    h_length;         /* length of address */  
    char **h_addr_list;      /* list of addresses */  
}
```

# Address Lookup

---

Known hosts

Known protocols (IP, TCP, UDP, ...)

Known services (TELNET, FTP, WWW, ...)



# Address Lookup – Known Hosts (Cont'd)

---

Two functions related to `gethostent()`

`sethostent`

- `void sethostent(int stayopen);`
- Open host database if it is not already open
- Rewind it if it is already open

`endhostent`

- `void endhostent(void);`
- Close the host database

NOTE: The two functions have different meanings when lookup via DNS

# Address Lookup, an Example

## – Get All Known Hosts

---

```
int main() {
    int i;
    char buf[64];
    struct hostent *h;
    while((h = gethostent()) != NULL) {
        if(h->h_addrtype != AF_INET)
            continue;
        printf("name=%s, addr={ ", h->h_name);
        for(i = 0; h->h_addr_list[i] != NULL; i++) {
            printf("%s ", inet_ntop(AF_INET,
                                    h->h_addr_list[i], buf, sizeof(buf)));
        }
        printf("}\n");
    }
    return(0);
}
```

```
name=cshome.cs.nctu.edu.tw, addr={ 140.113.235.101 }
name=csduty.cs.nctu.edu.tw, addr={ 140.113.235.102 }
```

# Address Lookup

## – Known Protocols

---

How do we get all known protocols in the system?

Check the `/etc/protocols` file

In a UNIX program, we can get all known protocols by the function:

- `struct protoent *getprotoent(void);`
- Returns: valid pointer if success, or NULL on error
- NOTE: `getprotoent()` is not thread-safe

```
struct protoent {  
    char    *p_name;           /* official protocol name */  
    char    **p_aliases;       /* alias list */  
    int      p_proto;          /* protocol number */  
}
```

# Address Lookup

## – Known Protocols (Cont'd)

---

Two functions related to `getprotoent()`

`setprotoent`

- `void setprotoent(int stayopen);`
- Open protocol database if it is not already open
- Rewind it if it is already opened

`endprotoent`

- `void endprotoent(void);`
- Close the protocol database

# Address Lookup, an Example

## – Get All Known Protocols

---

```
int main() {
    int i;
    struct protoent *p;
    while((p = getprotoent()) != NULL) {
        printf("name=%s (%d), ", p->p_name, p->p_proto);
        printf("alias={ ");
        for(i = 0; p->p_aliases[i] != NULL; i++)
            printf("%s ", p->p_aliases[i]);
        printf("}\n");
    }
    return(0);
}
```

```
name=ip (0), alias={ IP }
name=icmp (1), alias={ ICMP }
name=igmp (2), alias={ IGMP }
...
name=tcp (6), alias={ TCP }
...
name=udp (17), alias={ UDP }
...
```

# Address Lookup

## – Get a Specific Protocol

---

In addition to iteratively list all protocols, we can retrieve information about a given protocol

- `struct protoent *getprotobyname(const char *name);`
- `struct protoent *getprotobynumber(int proto);`
- Returns: valid pointer if success, or NULL on error
- NOTE: Both the two functions are not thread-safe

### Behavior of `setprotoent`

- By default, `getprotocolby*` functions close the protocol database after a query
- If we use `setprotoent` with `stayopen = true (1)`, the `getprotocolby*` functions do not close the database

# Address Lookup

## – Known Services

---

How do we get all known services in the system?

Check the `/etc/services` file

In a UNIX program, we can get all known services by the function:

- `struct servent *getservent(void);`
- Returns: valid pointer if success, or NULL on error
- NOTE: `getservent()` is not thread-safe

```
struct servent {  
    char    *s_name;           /* official service name */  
    char    **s_aliases;       /* alias list */  
    int      s_port;           /* port number */  
    char    *s_proto;          /* protocol to use */  
}
```

# Address Lookup

## – Known Services (Cont'd)

---

Two functions related to getservent()

setservent

- `void setservent(int stayopen);`
- Open service database if it is not already open
- Rewind it if it is already open

endservent

- `void endprotoent(void);`
- Close the service database



# Address Lookup, an Example

## – Get All Known Services

---

```
int main() {
    int i;
    struct servent *s;
    while((s = getservent()) != NULL) {
        printf("name=%s (%s/%d), ",
            s->s_name, s->s_proto, ntohs(s->s_port));
        printf("alias={ ");
        for(i = 0; s->s_aliases[i] != NULL; i++)
            printf("%s ", s->s_aliases[i]);
        printf("}\n");
    }
    return(0);
}
```

```
...
name=telnet (tcp/23), alias={ }
name=ftp (tcp/21), alias={ }
name=pop3 (tcp/110), alias={ pop-3 }
name=www (tcp/80), alias={ http }
...
```

# Address Lookup

## – Get a Specific Service

---

In addition to iteratively list all services, we can retrieve information about a given service

- `struct servent *getservbyname(const char *name, const char *proto);`
- `struct servent *getservbyport(int port, const char *proto);`
  - NOTE: port should be in network byte order
- Returns: valid pointer if success, or NULL on error
- NOTE: Both the two functions are not thread-safe

### Behaviors of setservent

- By default, each `getservby*` function close the service database after a query
- If we use `setservent` with `stayopen = true (1)`, the `getservby*` function will not close the database

# Address Lookup, Hostname via DNS

---

In addition to iteratively list all known hosts from system database, we can retrieve information about a given host via DNS

- `struct hostent *gethostbyname(const char *name);`
- `struct hostent *gethostbyaddr(const void *addr, socklen_t len, int type);`
  - type can be either `AF_INET` or `AF_INET6`
- Returns: valid pointer if success, or `NULL` on error
- NOTE: Both the two functions are not thread-safe

## Behaviors of `sethostent`

- By default, `gethostby*` functions query DNS using UDP protocol
- If we use `sethostent` with `stayopen = true (1)`, the `gethostby*` functions use TCP to query the DNS (and keep the connection alive)

# Thread-Safe Query of Address and Port (1/5)

---

## The getaddrinfo function

- `int getaddrinfo(const char *node, const char *service, const struct addrinfo *hints, struct addrinfo **res);`

## Parameters

- node: the node to be queried (name or address)
- service: name of the service
- hints: query criteria
  - flags (see the next slide)
  - address family (AF\_INET/AF\_INET6)
  - socktype (SOCK\_DGRAM/SOCK\_STREAM, can be 0)
  - protocol (can be 0)
  - Other fields must be zero
- res: return the queried result

Returns: zero if success, or nonzero error code on error

# Thread-Safe Query of Address and Port (2/5)

---

The addrinfo data structure

```
struct addrinfo {  
    int             ai_flags;  
    int             ai_family;  
    int             ai_socktype;  
    int             ai_protocol;  
    size_t          ai_addrlen;  
    struct sockaddr *ai_addr;  
    char            *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

# Thread-Safe Query of Address and Port (3/5)

---

The ai\_flags

Flag	Description
AI_ADDRCONFIG	Query for whichever address type (IPv4 or IPv6) is configured
AI_ALL	Look for both IPv4 and IPv6 addresses (used only with AI_V4MAPPED)
AI_CANONNAME	Request a canonical name (as opposed to an alias)
AI_NUMERICHOST	Return the host address in numeric format
AI_NUMERICSERV	Return the service as a port number
AI_PASSIVE	Socket address is intended to be bound for listening
AI_V4MAPPED	If no IPv6 addresses are found, return IPv4 addresses mapped in IPv6 format

# Thread-Safe Query of Address and Port (4/5)

---

Handle error returned from getaddrinfo

If getaddrinfo fails, we can not use perror or strerror to generate an error message

We need to call gai\_strerror to convert the error code returned into an error message

- `const char *gai_strerror(int error);`

# Thread-Safe Query of Address and Port (5/5)

```
$ ./getaddrinfo google.com www
74.125.45.100:80
209.85.171.100:80
74.125.67.100:80
```

```
int main(int argc, char *argv[]) {
    int s;
    struct addrinfo hints, *result, *rp;
    if (argc < 3) {
        fprintf(stderr, "usage: %s host port\n", argv[0]);
        exit(-1);
    }
    bzero(&hints, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;          /* allow IPv4 or IPv6 */
    hints.ai_socktype = SOCK_STREAM; /* stream socket */
    hints.ai_flags = 0;
    hints.ai_protocol = 0;              /* any protocol */
    if((s=getaddrinfo(argv[1], argv[2], &hints, &result))!=0) {
        fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(s));
        exit(-1);
    }
    for(rp = result; rp != NULL; rp = rp->ai_next) {
        struct sockaddr_in *p = (struct sockaddr_in*) rp->ai_addr;
        printf( "%s:%d\n", inet_ntoa(p->sin_addr), ntohs(p->sin_port));
    }
    return(0);
}
```



# Thread-Safe Query of Name and Service

---

The inverse function of `getaddrinfo`

## Synopsis

- `int getnameinfo(const struct sockaddr *sa, socklen_t salen, char *host, size_t hostlen, char *serv, size_t servlen, int flags);`
- Returns: zero if success, or nonzero on error

Flag	Description
NI_DGRAM	The service is datagram (UDP) based rather than stream (TCP)
NI_NAMEREQD	An error is returned if the hostname cannot be determined
NI_NOFQDN	Return only the hostname part of the fully qualified domain name for local hosts
NI_NUMERICHOST	Return the numeric form of the host address instead of the name
NI_NUMERICSERV	Return the numeric form of the service address (i.e., the port number) instead of the name

# Thread-Safe Query of Name and Service, an Example

---

```
int main(int argc, char *argv[]) {
    struct sockaddr_in sin;
    char host[64], serv[64];
    int s;
    if (argc < 3) {
        fprintf(stderr, "usage: %s ip port\n", argv[0]);
        exit(-1);
    }
    bzero(&sin, sizeof(sin));
    sin.sin_family = AF_INET;
    sin.sin_addr.s_addr = inet_addr(argv[1]);
    sin.sin_port = htons(atoi(argv[2]));
    if((s = getnameinfo((struct sockaddr*)&sin, sizeof(sin),
        host, sizeof(host), serv, sizeof(serv), 0)) != 0) {
        fprintf(stderr, "getnameinfo: %s\n", gai_strerror(s));
        exit(-1);
    }
    printf("%s:%s\n", host, serv);
    return(0);
}
```

```
$ ./getnameinfo 74.125.45.100 80
yx-in-f100.google.com:www
```

# Associate Addresses with Sockets

---

Usually, a client does not need to bind an address with a socket

- The server automatically chooses the address for the socket

However, a server has to bind an address with a socket

- `int bind(int sockfd, const struct sockaddr *addr, socklen_t addrlen);`
- Returns: zero if success, or -1 on error
- The address we specify must be valid for the machine
  - It can be zero – bound to all interfaces
- The port number in the address cannot be less than 1,024 (only superuser can do that)
- Usually, only one socket endpoint can be bound to a given address

# Discover the Address Bound to a Socket

---

Get the local address bound to a socket

- `int getsockname(int s, struct sockaddr *name, socklen_t *namelen);`

Get the remote address bound to a socket

- If the socket is connected to a peer
- `int getpeername(int s, struct sockaddr *name, socklen_t *namelen);`

Notes

- The *name* and the *namelen* must be declared first
- Before calling `getsockname` or `getpeername`, the *namelen* must be set to the length of the *name* data structure

# Connection Establishment

---

If a client is dealing with a connection-oriented network service (SOCK\_STREAM)

It has to create a connection before exchanging data

- `int connect(int sockfd, const struct sockaddr *serv_addr, socklen_t addrlen);`
- Returns: zero if success, or -1 on error
- If *sockfd* is not bound to an address, connect will bind a default address for the caller

# Listen for an Incoming Connection

---

A server can announce that it is willing to accept connect requests

The *listen* function

- `int listen(int sockfd, int backlog);`
- Returns: zero if success, or -1 on error
- `backlog`
  - The number of outstanding connect requests in a queue
  - In Linux, the max allowable is 128 (defined by the SOMAXCONN constant)
  - Once the queue is full, the system will reject additional connect requests
  - The *backlog* must be chosen based on the expected load of the server

# Accept Incoming Connections

---

Once a server has called `listen`, the socket used can receive connect requests

The *accept* function

- `int accept(int sockfd, struct sockaddr *addr, socklen_t *len);`
- Returns: file (socket) descriptor if success, or -1 on error
- The returned descriptor is the socket connected to the client
- This new socket descriptor has the same socket type and address family as the original *sockfd*
- The *addr* holds the address and the port of the client
  - It can be NULL if we do not need these information
- If no connect requests are pending, `accept` will block until one arrives

# Establish IPv4 TCP Connections – Summary

---

## Server side

- `fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP)`
- Provide a `sockaddr_in` data structure – `sin`
- `bind(fd, (struct sockaddr*) &sin, sizeof(sin))`
- `listen(fd, backlog)`
- `pfd = accept(fd, &psin, sizeof(psin))`

## Client side

- `fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP)`
- Provide a `sockaddr_in` data structure – `sin`
- `connect(fd, (struct sockaddr*) &sin, sizeof(sin))`



# Data Transfer – Send Data

---

We can use the *write* function to send data via descriptors

However, there are more flexible functions can be used

- For *connection oriented* only

```
ssize_t send(int sockfd, const void *buf, size_t nbytes,  
             int flags);
```

- It is equivalent to *write* if the *flags* is set to zero

- For both *Connection oriented* and *connectionless*

```
ssize_t sendto(int sockfd, const void *buf, size_t nbytes,  
               int flags, const struct sockaddr *destaddr,  
               socklen_t destlen);
```

- We have to specify a *destaddr* in connectionless mode
- Returns: number of bytes sent if success, or -1 on error

# Data Transfer – Send Data (Cont'd)

---

## Flags

Flag	Description
MSG_DONTROUTE	Don't route packet outside of local network
MSG_DONTWAIT	Enable non-blocking operation (equivalent to using O_NONBLOCK)
MSG_EOR	This is the end of record if supported by protocol
MSG_OOB	Send out-of-band data if supported by protocol

# Data Transfer – Receive Data

---

We can use the *read* function to receive data via descriptors

There are also more flexible functions can be used

- For *connection oriented* only  
`ssize_t recv(int sockfd, void *buf, size_t nbytes, int flags);`
- It is equivalent to *read* if the *flags* is set to zero
- For both *Connection oriented* and *connectionless*  
`ssize_t recvfrom(int sockfd, void *buf, size_t len,  
int flags,  
struct sockaddr *addr, socklen_t *addrlen);`
- The *addr* contains the address/port of the data sender (for connectionless mode only)
- Returns: number of bytes received if success, 0 if no messages are available and peer has done an orderly shutdown, or -1 on error

# Data Transfer – Receive Data (Cont'd)

---

## Flags

Flag	Description
MSG_OOB	Receive out-of-band data if supported by protocol
MSG_PEEK	Return packet contents without consuming packet
MSG_TRUNC	Return that the real length of the packet, even if it was longer than the passed buffer (Only valid for packet sockets)
MSG_WAITALL	Wait until all data is available, i.e., the passed buffer is all filled (SOCK_STREAM only)

# A Server Example – TCP ECHO Server (1/5)

---

```
int main(int argc, char *argv[]) {
    pid_t pid;
    int fd, pfd, val;
    struct sockaddr_in sin, psin;
    if(argc < 2) {
        fprintf(stderr, "usage: %s port\n", argv[0]);
        return(-1);
    }
    signal(SIGCHLD, SIG_IGN);
    if((fd = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0) {
        perror("socket");
        return(-1);
    }
    val = 1;
    if(setsockopt(fd,
        SOL_SOCKET, SO_REUSEADDR, &val, sizeof(val)) < 0) {
        perror("setsockopt");
        return(-1);
    }
}
```

## A Server Example – TCP ECHO Server (2/5)

---

```
bzero(&sin, sizeof(sin));
sin.sin_family = AF_INET;
sin.sin_port = htons(atoi(argv[1]));
if(bind(fd, (struct sockaddr*) &sin, sizeof(sin)) < 0) {
    perror("bind");
    return(-1);
}
if(listen(fd, SOMAXCONN) < 0) {
    perror("listen");
    return(-1);
}
```

## A Server Example – TCP ECHO Server (3/5)

---

```
while(1) {
    val = sizeof(psin);
    bzero(&psin, sizeof(psin));
    if((pfd=accept(fd, (struct sockaddr*) &psin, &val))<0) {
        perror("accept");
        return(-1);
    }
    if((pid = fork()) < 0) {
        perror("fork");
        return(-1);
    } else if(pid == 0) {      /* child */
        close(fd);
        serv_client(pfd, &psin);
        exit(0);
    }
    /* parent */
    close(pfd);
}
```

## A Server Example – TCP ECHO Server (4/5)

---

```
void serv_client(int fd, struct sockaddr_in *sin) {
    int len;
    char buf[2048];
    printf("connected from %s:%d\n",
           inet_ntoa(sin->sin_addr), ntohs(sin->sin_port));
    while((len = recv(fd, buf, sizeof(buf), 0)) > 0) {
        if(send(fd, buf, len, 0) < 0) {
            perror("send");
            exit(-1);
        }
    }
    printf("disconnected from %s:%d\n",
           inet_ntoa(sin->sin_addr), ntohs(sin->sin_port));
    return;
}
```



# A Server Example – TCP ECHO Server (5/5)

---

## Running the server

- `$ ./echosrv 12345`
- If you run `netstat -na` command, you should see:
  - `tcp 0 0 0.0.0.0:12345 0.0.0.0:* LISTEN`

## Running the client (using telnet)

- Type something and press Enter
- You should see the same string echoed back
- Press `^]` and type quit to terminate the client

```
$ telnet localhost 12345
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.
test 123
test 123
hello world
hello world
^]
telnet> quit
Connection closed.
```

# Socket Options

---

The behavior of sockets can be controlled by options

## Interfaces

- `int setsockopt(int sockfd, int level, int option, const void *val, socklen_t len);`
- `int getsockopt(int sockfd, int level, int option, void *val, socklen_t *lenp);`
- Returns: zero if success, or -1 on error
- The *level* argument
  - Identify the protocol (by a protocol number) to apply
    - For example, IPPROTO\_IP, IPPROTO\_TCP, ...
  - If the option is a generic socket-level option, then level is set to SOL\_SOCKET

# Generic Socket Options

Option	Type of val	Description
SO_ACCEPTCONN	int	Return whether a socket is enabled for listening (getsockopt only)
SO_BROADCAST	int	Broadcast datagrams if *val is nonzero
SO_DEBUG	int	Debugging in network drivers enabled if *val is nonzero
SO_DONTROUTE	int	Bypass normal routing if *val is nonzero
SO_ERROR	int	Return and clear pending socket error (getsockopt only)
SO_KEEPALIVE	int	Periodic keep-alive messages enabled if *val is nonzero
SO_LINGER	struct linger	Delay time when unsent messages exist and socket is closed
SO_OOBINLINE	int	Out-of-band data placed inline with normal data if *val is nonzero

# Generic Socket Options (Cont'd)

Option	Type of val	Description
SO_RCVBUF	int	The size in bytes of the receive buffer
SO_RCVLOWAT	int	The minimum amount of data in bytes to return on a receive call
SO_RCVTIMEO	struct timeval	The timeout value for a socket receive call
SO_REUSEADDR	int	Reuse addresses in bind if *val is nonzero
SO_SNDBUF	int	The size in bytes of the send buffer
SO_SNDLOWAT	int	The minimum amount of data in bytes to transmit in a send call
SO_SNDTIMEO	struct timeval	The timeout value for a socket send call
SO_TYPE	int	Identify the socket type (getsockopt only)

# Q & A

---